

## Chapter 11

### Resource Access Protocols

What are you about to learn? .....	2
11.1 Scheduling Dependent Task Sets .....	3
11.2 The Priority Inheritance Protocol .....	4
11.3 The Priority Ceiling Protocol .....	7
11.4 Scheduling Anomalies in Dependent Task Sets .....	13
Points to Remember .....	18

# Objectives

---

## What are you about to learn?

### Knowledge Objectives

- Understand the problems that come with *dependent* task sets.
- Be able to explain the phenomenon of priority inversion.
- Understand the two resource access protocol “priority inheritance protocol” and “priority ceiling protocol”.
- Know about some scheduling anomalies in dependent task sets.

### Skill Objectives

- Ability to define priority ceilings for a given task set.
- Ability to draw the plan for a dependent task set using either the priority inheritance mechanism or priority ceiling protocol, and the RM algorithm.

## 11.1 Scheduling Dependent Task Sets


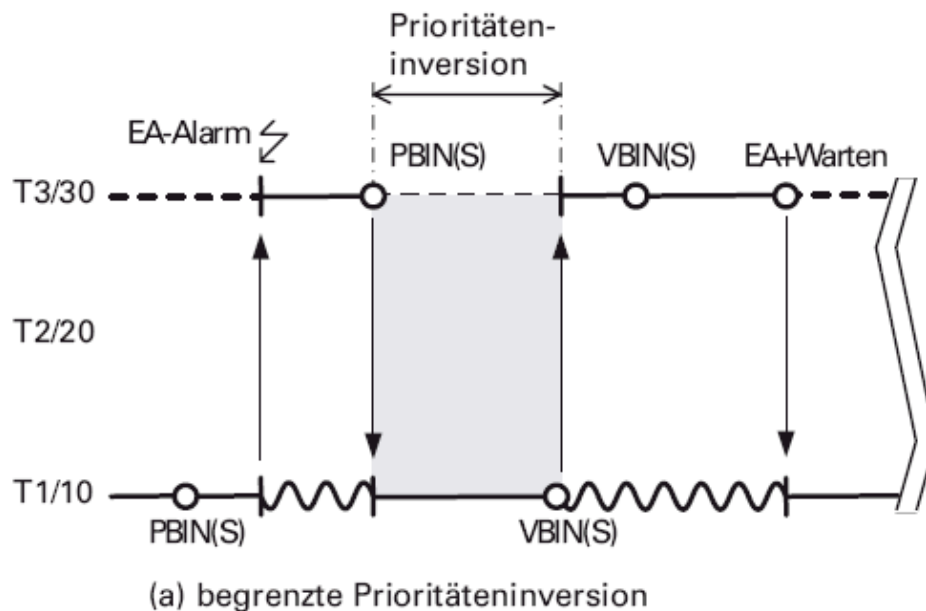
### 11.1 Scheduling Dependent Task Sets

Tasks become **dependent** on each other if they access the same resource (beside the processing unit itself).

Scheduling dependent task sets can introduce additional challenges. One problem that appears with priority based scheduling is **priority inversion**.

Priority inversion occurs when a low priority task T1 locks a resource that is then requested by a high priority task T3 (**bounded priority inversion**). T1 will block T3 via the shared resource, which is the same as if T1 has a higher priority than T3 (inversion of priority).

The priority inversion becomes **unbounded** when a third task T2 with a priority between those of T1 and T3 preempts T1. This is also called a **livelock**.



## 11.2 The Priority Inheritance Protocol

---


### 11.2 The Priority Inheritance Protocol

The priority inheritance protocol can be defined as follows:

- Jobs are scheduled based on their active priorities. Jobs with the same priority are executed on a first come first serve basis.
- When job  $J_i$  tries to enter a critical section  $z_{i,j}$  and resource  $R_{i,j}$  is already held by a lower priority job,  $J_i$  will be blocked.  $J_i$  is said to be blocked by the task that holds the resource. Otherwise,  $J_i$  enters the critical section  $z_{i,j}$ .
- When a job  $J_i$  is blocked on a semaphore, it transmits its active priority to the job, say  $J_k$ , that holds that semaphore. Hence,  $J_k$  resumes and executes the rest of its critical section with priority  $p_k = p_i$ .  $J_k$  is said to inherit the priority of  $J_i$ . In general, a task inherits the highest priority of the jobs blocked by it.
- When  $J_k$  exits a critical section, it unlocks the semaphore, and the highest-priority job, if any, blocked on that semaphore is awakened. Moreover, the active priority of  $J_k$  is updated as follows: if no other jobs are blocked by  $J_k$ ,  $p_k$  is set to its nominal priority  $P_k$ , otherwise it is set to the highest priority of the jobs blocked by  $J_k$ .
- Priority inheritance is transitive; that is if a job  $J_3$  blocks a job  $J_2$ , and  $J_2$  blocks a job  $J_1$ , then  $J_3$  inherits the priority of  $J_1$  via  $J_2$ .

See the following example with three jobs with priorities  $P_0 > P_1 > P_2$ .

## 11.2 The Priority Inheritance Protocol



## 11.2 The Priority Inheritance Protocol

Let  $B_i$  be the maximum blocking time, due to lower-priority jobs, that a job  $J_i$  may experience. Then a set of  $n$  periodic tasks using the priority inheritance protocol can be scheduled by the rate monotonic algorithm if

$$\forall i, 1 \leq i \leq n, \sum_{k=1}^i \frac{C_k}{T_k} + \frac{B_i}{T_i} \leq i(2^{1/i} - 1)$$

This can be simplified to a less tight condition

$$\sum_{i=1}^n \frac{C_i}{T_i} + \max\left(\frac{B_1}{T_1}, \dots, \frac{B_n}{T_n}\right) \leq n(2^{1/n} - 1)$$

Example:

	$C_i$	$T_i$	$B_i$
$J_1$	1	2	1
$J_2$	1	4	1
$J_3$	2	8	0

In this case the periods of the tasks are harmonic, the utilization bound for rate monotonic scheduling becomes 100%.

## 11.3 The Priority Ceiling Protocol

---

### 11.3 The Priority Ceiling Protocol

The priority ceiling protocol can be defined as follows:

- Each semaphore  $S_k$  is assigned a priority ceiling  $C(S_k)$  equal to the priority of the highest-priority job that can lock it.  $C(S_k)$  is a static value that can be computed off-line.
- Let  $J_i$  be the job with the highest priority among all jobs ready to run; thus  $J_i$  is assigned the processor.
- Let  $S^*$  be the semaphore with the highest priority ceiling among all the semaphores currently locked by jobs other than  $J_i$ , and let  $C(S^*)$  be its ceiling.
- To enter a critical section guarded by a semaphore  $S_k$ ,  $J_i$  must have a priority higher than  $C(S^*)$ . If  $P_i \leq C(S^*)$ , the lock on  $S_k$  is denied and  $J_i$  is said to be blocked on semaphore  $S^*$  by the job that holds the lock on  $S^*$ .
- When a job  $J_i$  is blocked on a semaphore, it transmits its priority to the job, say  $J_k$ , that holds that semaphore. Hence,  $J_k$  resumes and executes the rest of its critical section with the priority of  $J_i$ .  $J_k$  is said to inherit the priority of  $J_i$ .
- When  $J_k$  exits a critical section, it unlocks the semaphore and the highest-priority job, if any, blocked on that semaphore is awakened. Moreover, the active priority of  $J_k$  is updated as follows: if no other jobs are blocked by  $J_k$ ,  $p_k$  is set to the nominal priority  $P_k$ ; otherwise, it is set to the highest priority of the jobs blocked by  $J_k$ .
- Priority inheritance is transitive; that is if a job  $J_3$  blocks a job  $J_2$ , and  $J_2$  blocks a job  $J_1$ , then  $J_3$  inherits the priority of  $J_1$  via  $J_2$ .

## 11.3 The Priority Ceiling Protocol

---

See the following example with three jobs with priorities  $P_0 > P_1 > P_2$ .

- Job  $J_0$  sequentially accesses two critical sections guarded by semaphores  $S_0$  and  $S_1$
- Job  $J_1$  accesses only a critical section guarded by semaphore  $S_2$
- Job  $J_2$  uses semaphore  $S_2$  and then makes a nested access to  $S_1$ .

The semaphores are thus assigned the following priority ceilings:

- $C(S_0) = P_0$ , since  $J_0$  is the highest-priority task that tries to lock  $S_0$
- $C(S_1) = P_0$ , since  $J_0$  is the highest-priority task that tries to lock  $S_1$
- $C(S_2) = P_1$ , since  $J_1$  is the highest-priority task that tries to lock  $S_2$


The activation times of the three jobs are as follows:

- $J_0$  is activated at  $t = 8$
- $J_1$  is activated at  $t = 3$
- $J_2$  is activated at  $t = 0$

See the following diagram for the scenario.



## 11.3 The Priority Ceiling Protocol



## 11.3 The Priority Ceiling Protocol

---

Scenario description:

- At time 0,  $J_2$  is activated, and since it is the only job ready to run, it starts executing and later locks semaphore  $S_2$ .
- At time 3,  $J_1$  becomes ready and preempts  $J_2$ .
- At time 4,  $J_1$  attempts to lock  $S_2$ , but it is blocked by the protocol because  $P_1$  is not greater than  $C(S_2)$ . Then,  $J_2$  inherits the priority of  $J_1$  and resumes its execution.
- At time 6,  $J_2$  successfully enters its nested critical section by locking  $S_1$ . Note that  $J_2$  is allowed to lock  $S_1$  because no semaphores are locked by other jobs.
- At time 8, while  $J_2$  is executing at a priority  $p_2 = P_1$ ,  $J_0$  becomes ready and preempts  $J_2$  because  $P_0 > p_2$ .
- At time 10,  $J_0$  attempts to lock  $S_0$ , which is not locked by any job. However,  $J_0$  is blocked by the protocol because its priority is not higher than  $C(S_1)$ , which is the highest ceiling among all semaphores currently locked by the other jobs. Since  $S_1$  is locked by  $J_2$ ,  $J_2$  inherits the priority of  $J_0$  and resumes its execution.
- At time 12,  $J_2$  exits its nested critical section, unlocks  $S_1$ , and, since  $J_0$  is awakened,  $J_2$  returns to priority  $p_2 = P_1$ . At this point,  $P_0 > C(S_2)$ ; hence,  $J_0$  preempts  $J_2$  and executes until completion.
- At time 15,  $J_0$  is completed, and  $J_2$  resumes its execution at a priority  $p_2 = P_1$ .
- At time 16,  $J_2$  exits its outer critical section, unlocks  $S_2$ , and, since  $J_1$  is awakened,  $J_2$  returns to its nominal priority  $P_2$ . At this point,  $J_1$  preempts  $J_2$  and executes until completion.

## 11.3 The Priority Ceiling Protocol

---

- At time 19,  $J_1$  is completed; thus  $J_2$  resumes its execution.

### Blocking time computation of the priority ceiling protocol:

The maximum blocking time of  $B_i$  of a job  $J_i$  can be computed as the duration of the longest critical section among those belonging to tasks with priority lower than  $P_i$  and guarded by a semaphore with ceiling higher than or equal to  $P_i$ . If  $D_{j,k}$  denotes the duration of the longest critical section of task  $\tau_j$  among those guarded by semaphore  $S_k$ , we can write

$$B_i = \max_{j,k} \{D_{j,k} \mid P_j < P_i, C(S_k) \geq P_i\}$$

Example: for each job  $J_i$ , the duration of the longest critical section among those guarded by the semaphore  $S_k$  is denoted by  $D_{i,k}$  and is stored in a table.  $D_{i,k} = 0$  means that job  $J_i$  does not use semaphore  $S_k$ . Semaphore ceilings are in parentheses:

## 11.3 The Priority Ceiling Protocol

---

	$S_1 (P_1)$	$S_2 (P_1)$	$S_3 (P_2)$
$J_1$	1	2	0
$J_2$	0	9	3
$J_3$	8	7	0
$J_4$	6	5	4

The tasks' blocking factors are computed as follows

$$\left\{ \begin{array}{l} B_1 = \max(8, 6, 9, 7, 5) = 9 \\ B_2 = \max(8, 6, 7, 5, 4) = 8 \\ B_3 = \max(6, 5, 4) = 6 \\ B_4 = 0 \end{array} \right.$$


## 11.4 Scheduling Anomalies in Dependent Task Sets

### 11.4 Scheduling Anomalies in Dependent Task Sets

When dealing with task sets with precedence relations executed in a multiprocessor environment we encounter some scheduling anomalies:

- adding resources (such as a processor) can make things worse
- relaxing constraints such as less precedence between tasks or lower execution time requirements can make things worse

This is illustrated in the following example with nine tasks and the following precedence relationships (values in parentheses are execution times):




$$\text{priority}(J_i) > \text{priority}(J_j) \quad \forall i < j$$


*values in parantheses are execution times*

## 11.4 Scheduling Anomalies in Dependent Task Sets

First case: optimal schedule on three processors  $P_1, P_2, P_3$ . Global completion time is 12.




Second case: [we add a processor](#)  $P_4$ . Tasks are allocated to the first available processor. Global completion time has increased to 15.



Increase to four processors  
Allocation to the first available processor


## 11.4 Scheduling Anomalies in Dependent Task Sets

Third case: [we reduce all computation times by one time unit](#). Tasks are allocated to the first available processor. Global completion time has increased to 13.



Computation time reduced by one time unit


## 11.4 Scheduling Anomalies in Dependent Task Sets



Fourth case: **we remove some precedence constraints**. Global completion time has increased to 16.

$$\text{priority}(J_i) > \text{priority}(J_j) \quad \forall i < j$$

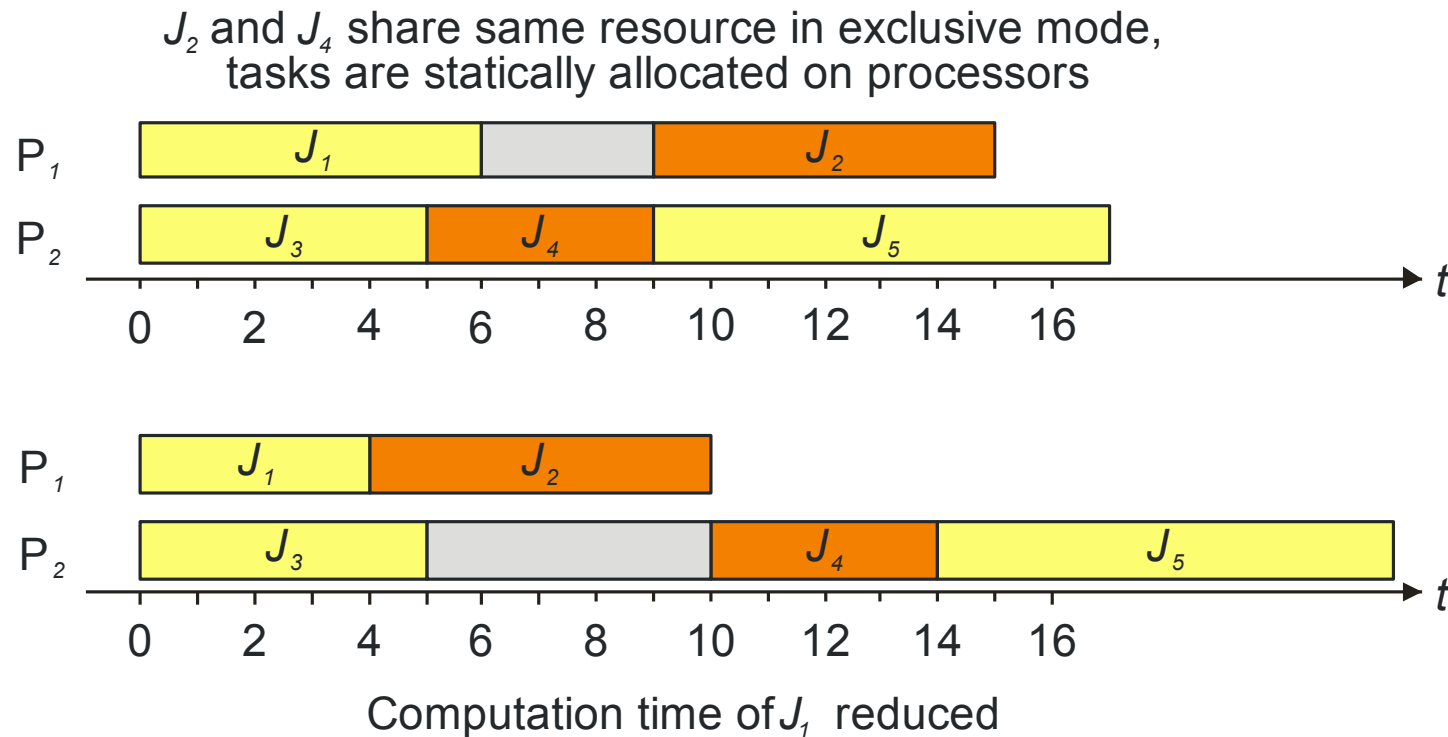
values in parentheses are execution times





## 11.4 Scheduling Anomalies in Dependent Task Sets

Fifth case: [we share a resource in exclusive mode, and reduce computation time](#). Global completion time increases from 16 to 22.



## Points to Remember

---

## Points to Remember